QUIZ • 10 MIN

Machine Learning System Design

Submit your assignment

DUE Aug 16, 11:59 PM PDT **ATTEMPTS** 3 every 8 hours

Resume

Receive grade

TO PASS 80% or higher

Grade 80%

View Feedback We keep your highest score

Machine Learning System Design

TOTAL POINTS 5

1. You are working on a spam classification system using regularized logistic regression. "Spam" is a positive class (y = 1) and "not spam" is the negative class (y = 0). You have trained your classifier and there are m = 1000 examples in the cross-validation set. The chart of predicted class vs. actual class is: 1 point

	Actual Class: 1	Actual Class: 0
Predicted Class: 1	85	890
Predicted Class: 0	15	10

For reference:

- Accuracy = (true positives + true negatives) / (total examples)
- Precision = (true positives) / (true positives + false positives)
- Recall = (true positives) / (true positives + false negatives)
- F_1 score = (2 * precision * recall) / (precision + recall)

What is the classifier's precision (as a value from 0 to 1)?

Enter your answer in the box below. If necessary, provide at least two values after the decimal point.

Enter answer here

2. Suppose a massive dataset is available for training a learning algorithm. Training on a lot of data is likely to give good performance when two of the following conditions hold true. 1 point

Which are the two?

- [ ] Our learning algorithm is able to represent fairly complex functions (for example, if we train a neural network or other model with a large number of parameters).
- [ ] When we are willing to include high order polynomial features of x (such as x_1^2, x_2^2, x_1x_2, etc.).
- [ ] The classes are not too skewed.
- [ ] A human expert on the application domain can confidently predict y when given only the features x (or more generally, if we have some way to be confident that x contains sufficient information to predict y accurately).

3. Suppose you have trained a logistic regression classifier which is outputing $h_\theta(x)$. 1 point

Currently, you predict 1 if $h_\theta(x) \geq \text{threshold}$, and predict 0 if $h_\theta(x) lt \text{threshold}$, where currently the threshold is set to 0.5.

Suppose you **decrease** the threshold to 0.1. Which of the following are true? Check all that apply.

- [ ] The classifier is likely to have unchanged precision and recall, but lower accuracy.
- [ ] The classifier is likely to now have higher recall.
- [ ] The classifier is likely to now have higher precision.
- [ ] The classifier is likely to have unchanged precision and recall, but higher accuracy.

4. Suppose you are working on a spam classifier, where spam emails are positive examples ($y = 1$) and non-spam emails are negative examples ($y = 0$). You have a training set of emails in which 99% of the emails are non-spam and the other 1% is spam. Which of the following statements are true? Check all that apply. 1 point

- [ ] If you always predict non-spam (output $y = 0$), your classifier will have an accuracy of 99%.
- [ ] If you always predict non-spam (output $y = 0$), your classifier will have 99% accuracy on the training set, and it will likely perform similarly on the cross validation set.
- [ ] A good classifier should have both a high precision and high recall on the cross validation set.
- [ ] If you always predict non-spam (output $y = 0$), your classifier will have 99% accuracy on the training set, but it will do much worse on the cross validation set because it has overfit the training data.

5. Which of the following statements are true? Check all that apply. 1 point

- [ ] After training a logistic regression classifier, you **must** use 0.5 as your threshold for predicting whether an example is positive or negative.
- [ ] Using a **very large** training set makes it unlikely for model to overfit the training data.
- [ ] It is a good idea to spend a lot of time collecting a **large** amount of data before building your first version of a learning algorithm.
- [ ] If your model is underfitting the training set, then obtaining more data is likely to help.
- [ ] The "error analysis" process of manually examining the examples which your algorithm got wrong can help suggest what are good steps to take (e.g., developing new features) to improve your algorithm's performance.

- [ ] I understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.

Learn more about Coursera's Honor Code

Save Submit